

Systems and Information Security SEGSI

Topic 7
Architectures

Architectures

- ▶ When designing a secure architecture some principles should be considered and, if possible, implemented
- ▶ A security failure might have has a starting point
 - ▶ A failure of a asset
 - ▶ A failure of an application
 - ▶ An exploit of a vulnerability
- ▶ The consequences might be an erroneous state or the disclosure of unintended information


Architectures

- ▶ In this topic we will focus on asset and application failures
- ▶ When designing an infrastructure we already saw (Topic 5) that failures might occur on any asset, and the strategies that might be implemented
 - ▶ Preventive fault detection
 - ▶ Retrospective fault detection
- ▶ Whatever strategy is implemented, the considerations for hardware and software differ

Architectures

- ▶ In hardware terms, a fault-tolerant hardware can be based on triple modular redundancy (TMR)
 - ▶ There are three identical and replicated components that receive the same input and whose outputs are compared
 - ▶ If an output is different it is ignored and the respective component is assumed to be defective

Architectures



Source: Engineers Community

Note: alternative implementations might be used

Architectures

- ▶ The output comparator is generally a relatively simple hardware device
- ▶ Compare your entries and reject one that is different from the rest
- ▶ The selection of the “correct” exit is made through a vote in which the one with the most votes wins
- ▶ TMR's success in providing fault tolerance is based on two fundamental assumptions
 - ▶ Hardware components have no common design flaws
 - ▶ These components fail randomly and there is a low probability of simultaneous failure

Architectures

- ▶ TMR-related problems can be perceived under different perspectives
 - ▶ Problem 1
 - ▶ Is it clear that if one component outputs a different (and, as so, minority) result it is undoubtedly the incorrect one?
 - ▶ Problem 2
 - ▶ Only one voter?
 - ▶ It might be a SPOF
 - ▶ Problem 3
 - ▶ If one component is turned off and the remaining two outputs different results, what would be the option of the voter?

Architectures

- ▶ The solution to the above problems involves different approaches to architectural design
- ▶ Obviously, duplicating the architecture more and more is not a solution
- ▶ However, the assumptions stated on [slide 6](#) should / must be taken into consideration
- ▶ In addition, and if possible
 - ▶ Ensure controls that allow better granularity of the definition of “correct”
 - ▶ Consider implementing a double line of components or at least voters
 - ▶ Analyze and test procedures in the event of a tie

Architectures

- ▶ On slide 6 the following assumptions were presented
 - ▶ *Hardware components have no common design flaws*
 - ▶ *These components fail randomly and there is a low probability of simultaneous failure*
- ▶ However, it's not usually easy to hardware components to have no common design flaws
- ▶ If they all have the same architecture, a common error might state up, simultaneously
 - ▶ As well as on an application
- ▶ That is the thinking that drives redundancy and diversity

Architectures

- ▶ Redundancy is achieved if an environment is duplicated (at least), resulting in distinct and absolutely identical environments
- ▶ Diversity is achieved if different versions of an environment is deployed
- ▶ An easy way to differ them is to imagine systems implementation
 - ▶ Redundancy
 - ▶ Two or more systems with same operating system, programming languages, applications, etc.
 - ▶ Diversity
 - ▶ Two or more systems with different operating systems and / or programming languages and / or applications, etc.
- ▶ This idea can be improved by having different manufacturers or equivalent

Architectures

- ▶ The assumptions of TMR are not viable in software!
 - ▶ It is not recommended to simply replicate the same component, as this would cause them to have common design flaws
 - ▶ Therefore, the simultaneous failure of software components is virtually inevitable
 - ▶ Software systems must be diversified
- ▶ Thus, TMR analogies can be used on software

Architectures

- ▶ Software design should / must use diversity
 - ▶ Different versions of the system are designed and implemented differently
 - ▶ The goal is to have different failure modes!
 - ▶ As well as different approaches to design
 - ▶ Different paradigms - object oriented, functional, imperative, etc.
 - ▶ Different implementation languages
 - ▶ Different tools and development environments
 - ▶ Different algorithms in the implementation

Architectures

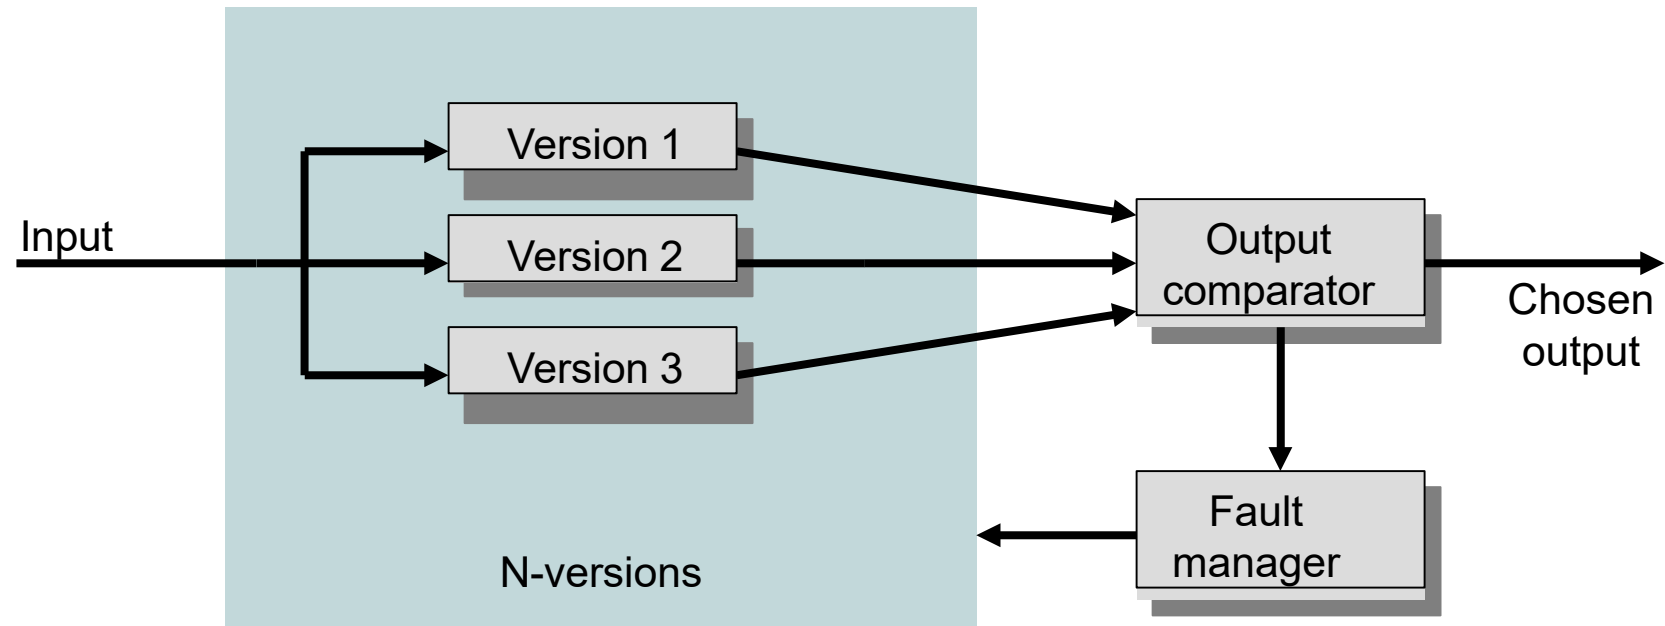
- ▶ Analogies of software TMR
 - ▶ Programming with N-versions
 - ▶ The same specification is implemented in different ways by different teams
 - ▶ All versions count simultaneously and the majority vote determines the chosen output
 - ▶ Most used approach in many models of commercial Airbus or Boeing aircraft and others
 - ▶ Recovery blocks
 - ▶ Several different versions of the same specification are designed and executed sequentially
 - ▶ An acceptance test is used to select the output that will be transmitted to the rest of the system

Architectures

- ▶ N-Versions
 - ▶ All versions receive the same input at the same time
 - ▶ Output comparator might have a maximum wait time for all versions to provide an output, specially for real-time systems
 - ▶ The most voted output is chose to continue for the next step / process
 - ▶ The most demanding paper here is the output comparator, in spite of being a simple software component that uses a voting mechanism to choose the final output
 - ▶ As with hardware-based systems, the output comparator is a simple software component that uses a voting mechanism to choose the final output

Architectures

► N-Versions

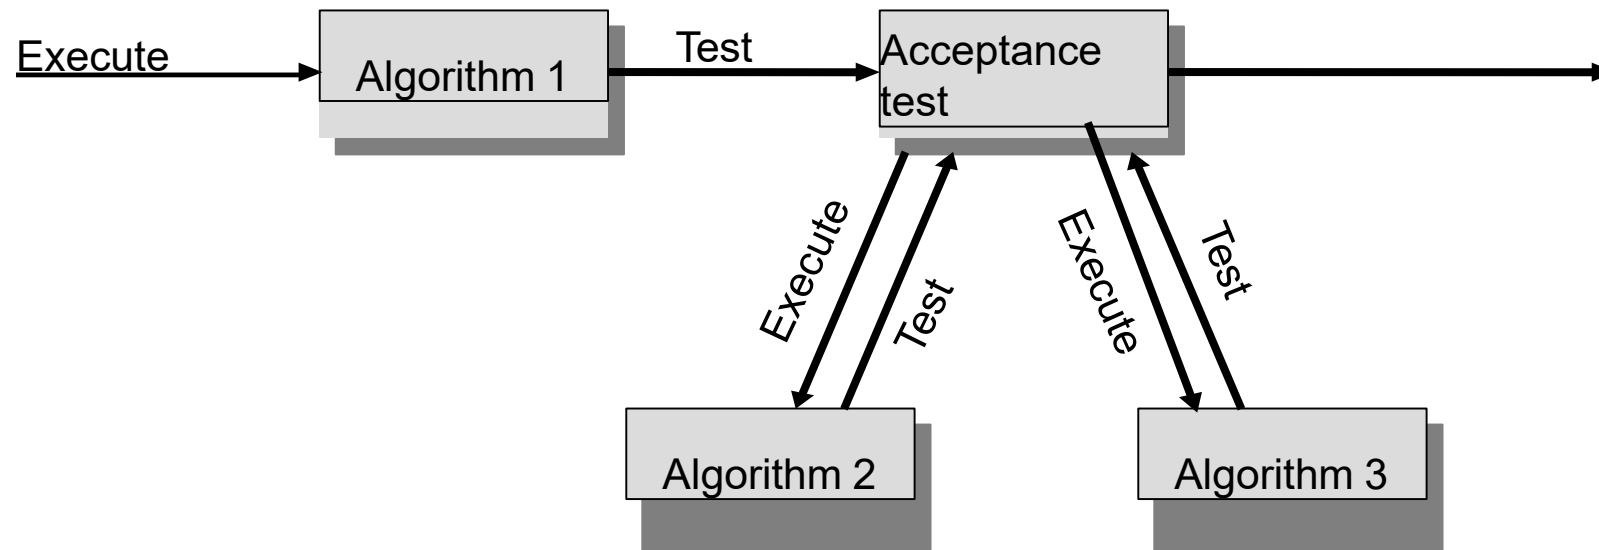


Architectures

- ▶ Recovery blocks
 - ▶ There are different architectures to assure same functionality
 - ▶ Each recovery block contains
 - ▶ 1. An algorithm
 - ▶ 2. An acceptance test (might be common for all algorithms)
 - ▶ Algorithm 2 will only be called if the test result from algorithm 1 is error (or to check algorithm 1 result), and same applies to call algorithm 3
 - ▶ It is mandatory that a different algorithm be executed on the various blocks in order to reduce the probabilities of common errors
 - ▶ The design of the acceptance test is complicated as it must be independent of the computation performed
 - ▶ There are complex problems with this approach in the case of real-time systems due to the sequential execution of the different redundant blocks

Architectures

► Recovery blocks



Architectures

- ▶ Whatever analogy is used, the perspectives of [slide 7](#) still applies
 - ▶ Problem 1
 - ▶ Is it clear that if one component (version or algorithm in this case) outputs a different (and, as so, minority) result it is undoubtedly the incorrect one?
 - ▶ Problem 2
 - ▶ Only one output comparator / acceptance test?
 - ▶ In spite of every algorithm can contain its own acceptance test on Recovery Blocks option
 - ▶ It might be a SPOF
 - ▶ Problem 3
 - ▶ If one version or algorithm is turned off and the remaining two outputs different results, what would be the option of the output comparator / acceptance test?